


LegalEase — Master System Flowchart (Dark)

End-to-End: User -> Frontend -> API (Browser) -> Backend -> AI / Modes -> Persistence -> Learning



Process Relationships — Quick Reference

CHAT uses MATTERS for scope validation (matter_policy) and FAISS matter index for retrieval.
DOCUMENTS feed FAISS (global or matter path) – same index CHAT queries in KB mode.
LEARNING reads chat interactions; writes boosts – never injects legal substance from Gemini into KB.
OPEN LAW / HYBRID use Gemini; KB mode uses Ollama only (GEMINI_KB_SYNTHESIS=0).
MATTER delete archives by default; documents unlinked from matter index on archive.
Frontend streamChat -> chat/stream -> chat_service -> mode_router -> (rag | web_intelligence).
Upload -> documents endpoint -> OCR/extract -> chunk/embed -> index_jobs -> FAISS + SQLite metadata.
Settings coach + neural tuning run in background; triggered by feedback, not during chat turn.
SaaS modules (billing, CRM, discovery) share SQLite user scope; independent API routers.